Use these AI tools if you want the ideal resume.

1. **Resumaker.ai** – **Resume & Cover letter generator**
2. **Resume.io** – **AI powered resume Builder**
3. **Jobhunt.com** – **AI powered job application assistant**
4. **Zety.com** – **Resume checker;Review and score your Resume**
5. **Jobscan.co** – **optimize your resume to get more interviews**
6. **Jobprofile.io** – **create a winning resume in minutes**
7. **Practiceinterviews.com** – **Chatbot bot interview preparation**
8. **Jobinterview.coach** – **The only complete AI interview coaching platform**
9. **App.yoodli.ai** – **improve your communication skill using AI**
10. **Carriercircles.com** – **Helps people affected by layoffs to bounce back**

AI tools for your daily life

1. **podcatstle.ai**
   - **Text to audio file**
   - **Audio to text file**

2. **Google translate**
   - **Translate document**
   - **Translate image**
   - **Translate website**

3. **humata.ai**
   - **Summarize the document**
   - **Clear the doubts in the document**

4. **ChatGPT**
   - **Give answer whatever you want till sept 2021**

5. **Google Gemini**
   - **Same like chatGPT upto current**

6. **Tome.app / Gamma**

   – **create a cartoon image**

   – **Document to PPT**

   – **Create PPT**

7. **Remove.bg**

   – **Remove background from image**

8. **fliki.ai** – **Create instant videos**

9. **fliki.ai** – **improve your communication skill using AI**

10. **Carriercircles.com** – **Helps people affected by layoffs to bounce back**

AI Interview Questions

1. **What is artificial intelligence?**

    Artificial intelligence (AI) refers to the simulation of human intelligence in machines that are programmed to think and act like humans.

2. **What are the different types of AI?**

    AI can be categorized into three types: narrow AI (also known as weak AI), general AI (also known as strong AI), and artificial superintelligence.

3. **What is machine learning?**

    Machine learning is a subset of AI that enables machines to learn from data and improve their performance over time without being explicitly programmed.

4. **Explain supervised learning.**

    Supervised learning is a type of machine learning where the model is trained on labeled data, and it learns to make predictions by mapping input data to output labels.

5. **What is unsupervised learning?**

    Unsupervised learning is a type of machine learning where the model is trained on unlabeled data, and it learns to find patterns or structures in the data without explicit guidance.

6. **Differentiate between classification and regression.**

    Classification is a type of supervised learning where the output variable is a category, while regression is a type of supervised learning where the output variable is a continuous value.

7. **What is reinforcement learning?**

Reinforcement learning is a type of machine learning where an agent learns to make decisions by interacting with an environment to maximize cumulative rewards.

8. Explain the bias-variance tradeoff.

   The bias-variance tradeoff is a fundamental concept in machine learning, which refers to the balance between a model's ability to capture the underlying patterns in the data (bias) and its sensitivity to variations in the data (variance).

9. What is overfitting, and how can it be prevented?

   Overfitting occurs when a model learns the training data too well, capturing noise or random fluctuations, leading to poor performance on unseen data. It can be prevented by using techniques like cross-validation, regularization, or collecting more data.

10. What is underfitting?

    Underfitting occurs when a model is too simple to capture the underlying structure of the data, resulting in poor performance both on the training and unseen data.

11. What are some common machine learning algorithms?

    Common machine learning algorithms include linear regression, logistic regression, decision trees, random forests, support vector machines, k-nearest neighbors, naive Bayes, neural networks, etc.

12. What is deep learning?

    Deep learning is a subset of machine learning that utilizes artificial neural networks with multiple layers (deep architectures) to learn complex patterns from large amounts of data.

13. Explain backpropagation.

    Backpropagation is a technique used to train neural networks by calculating the gradient of the loss function with respect to the weights of the network and adjusting the weights using gradient descent optimization.

14. What is a neural network?

A neural network is a computational model inspired by the structure and function of the human brain, consisting of interconnected nodes (neurons) organized in layers. It is capable of learning complex patterns and relationships from data.

15. What are activation functions, and why are they important in neural networks?

Activation functions introduce non-linearity to the output of a neuron in a neural network, allowing it to learn complex patterns and relationships in the data. Common activation functions include sigmoid, tanh, ReLU, Leaky ReLU, etc.

16. What is a convolutional neural network (CNN)?

A convolutional neural network (CNN) is a type of neural network that is particularly well-suited for tasks involving image recognition and classification. It uses convolutional layers to automatically learn hierarchical patterns and features from images.

17. What is a recurrent neural network (RNN)?

A recurrent neural network (RNN) is a type of neural network designed to handle sequential data by maintaining internal memory. It is commonly used in tasks such as natural language processing, speech recognition, and time series prediction.

18. Explain the vanishing gradient problem.

The vanishing gradient problem occurs during the training of deep neural networks when the gradients of the loss function become extremely small as they propagate backward through the network layers, leading to slow or stalled learning.

19. How can the vanishing gradient problem be mitigated in neural networks?

The vanishing gradient problem can be mitigated by using activation functions like ReLU, initializing the weights appropriately, using batch normalization, employing techniques like residual connections, or using alternative architectures like LSTM or GRU in RNNs.

20. What is transfer learning?

Transfer learning is a machine learning technique where a pre-trained model is used as a starting point for a new task, and then fine-tuned on a smaller dataset specific to the new task. It can help improve performance, especially when limited data is available.

21. Explain the concept of ensemble learning.

Ensemble learning is a machine learning technique where multiple models (learners) are combined to improve overall performance. Common methods include bagging, boosting, and stacking.

22. What is bagging?

Bagging (Bootstrap Aggregating) is an ensemble learning technique where multiple models are trained on different subsets of the training data with replacement, and their predictions are combined through averaging or voting to make the final prediction.

23. What is boosting?

Boosting is an ensemble learning technique where multiple weak learners are combined sequentially to create a strong learner. Each subsequent model focuses on the examples that the previous ones misclassified.

24. What is gradient boosting?

Gradient boosting is a popular boosting technique where weak learners (usually decision trees) are added sequentially to the ensemble, with each new learner trained to correct the errors made by the existing ensemble.

25. What is XGBoost, and why is it widely used?

XGBoost (Extreme Gradient Boosting) is an optimized implementation of gradient boosting, known for its speed and performance. It is widely used in machine learning competitions and various applications due to its efficiency and effectiveness.

26. What is a decision tree?

A decision tree is a tree-like model used for both classification and regression tasks. It breaks down a dataset into smaller subsets based on different attributes, leading to a tree-like decision structure.

27. How does a decision tree decide where to split?

A decision tree decides where to split based on criteria such as Gini impurity or entropy for classification tasks and mean squared error for regression tasks. It chooses the split that maximizes the homogeneity of the target variable in the resulting subsets.

28. What is Gini impurity?

Gini impurity is a measure of the impurity or disorder in a set of elements. In the context of decision trees, it is used as a criterion for splitting nodes, aiming to minimize impurity in the resulting child nodes.

29. What is entropy?

Entropy is a measure of randomness or uncertainty in a set of elements. In decision trees, it is used as a criterion for splitting nodes, aiming to maximize the information gain or decrease in entropy in the resulting child nodes.

30. What is hyperparameter tuning?

Hyperparameter tuning refers to the process of finding the optimal values for the hyperparameters of a machine learning model. It involves techniques such as grid search, random search, or Bayesian optimization to search the hyperparameter space efficiently.

31. What are hyperparameters?

Hyperparameters are parameters of a machine learning model that are set before the training process begins and remain constant during training. Examples include learning rate, regularization strength, number of hidden layers, etc.

32. What is grid search?

Grid search is a hyperparameter tuning technique where a grid of hyperparameter values is defined, and the model is trained and evaluated for each combination of hyperparameters. The combination that yields the best performance is selected.

33. What is random search?

Random search is a hyperparameter tuning technique where hyperparameters are sampled randomly from predefined distributions, and the model is trained and evaluated for each random combination of hyperparameters. It is often more efficient than grid search.

34. Explain the concept of bias in machine learning.

Bias in machine learning refers to the error introduced by overly simplistic assumptions in the learning algorithm, leading to underfitting and poor performance on both training and unseen data.

35. Explain the concept of variance in machine learning.

Variance in machine learning refers to the error introduced by the algorithm's sensitivity to fluctuations in the training data, leading to overfitting and poor generalization to unseen data.

36. What is the tradeoff between bias and variance?

The bias-variance tradeoff refers to the balancing act between minimizing bias and variance in a machine learning model. Decreasing bias often increases variance, and vice versa. The goal is to find the right balance to achieve optimal model performance.

37. What is precision?

Precision is a metric used to evaluate the performance of a classification model, representing the ratio of true positive predictions to the total number of positive predictions made by the model.

38. What is recall?

Recall is a metric used to evaluate the performance of a classification model, representing the ratio of true positive predictions to the total number of actual positive instances in the data.

39. What is F1-score?

F1-score is the harmonic mean of precision and recall, providing a single metric that balances both measures. It is commonly used as an evaluation metric for classification models, especially when dealing with imbalanced datasets.

40. What is accuracy?

Accuracy is a metric used to evaluate the performance of a classification model, representing the ratio of correctly predicted instances to the total number of instances in the dataset.

41. What is ROC curve, and what does it represent?

ROC (Receiver Operating Characteristic) curve is a graphical plot that illustrates the performance of a binary classification model across different threshold settings. It shows the tradeoff between true positive rate (TPR) and false positive rate (FPR).

42. What is AUC-ROC?

AUC-ROC (Area Under the ROC Curve) is a metric used to quantify the overall performance of a binary classification model. It represents the area under the ROC curve, with a higher value indicating better discrimination between positive and negative instances.

43. What is cross-validation?

Cross-validation is a technique used to assess the performance and generalization ability of a machine learning model. It involves splitting the dataset into multiple subsets, training the model on some subsets, and evaluating it on the remaining subset in an iterative manner.

44. Explain the K-fold cross-validation technique.

K-fold cross-validation is a cross-validation technique where the dataset is divided into K equal-sized folds. The model is trained K times, each time using K-1 folds for training and the remaining fold for validation. The final performance metric is the average of the K validation results.

45. What is the curse of dimensionality?

The curse of dimensionality refers to the phenomena where the performance of machine learning algorithms deteriorates as the number of features or dimensions in the data increases, leading to increased computational complexity and decreased generalization ability.

46. How can the curse of dimensionality be mitigated?

The curse of dimensionality can be mitigated by techniques such as feature selection, dimensionality reduction (e.g., PCA, t-SNE), using appropriate algorithms that are robust to high-dimensional data, or collecting more data to reduce sparsity.

47. What is feature engineering?

Feature engineering is the process of selecting, transforming, or creating new features from raw data to improve the performance of machine learning models. It involves domain knowledge, creativity, and experimentation.

48. What are some common feature engineering techniques?

Common feature engineering techniques include imputation of missing values, scaling and normalization, one-hot encoding for categorical variables, feature extraction from text or images, creating interaction or polynomial features, etc.

49. What is the importance of domain knowledge in feature engineering?

Domain knowledge is crucial in feature engineering as it helps identify relevant features, understand the relationships between features and the target variable, and guide the selection and transformation of features to improve model performance.

50. What is the role of bias in machine learning algorithms?

Bias in machine learning algorithms can result from inherent assumptions or limitations in the learning algorithm, leading to systematic errors in predictions. Addressing bias is crucial to ensure fairness, transparency, and accuracy in AI systems.

51. How can bias in machine learning algorithms be detected and mitigated?

Bias in machine learning algorithms can be detected and mitigated through various techniques such as careful selection and preprocessing of data, fairness-aware learning algorithms, bias audits, and diverse model evaluation.

52. What is fairness in machine learning, and why is it important?

Fairness in machine learning refers to the absence of discrimination or bias in the decisions made by AI systems across different demographic groups. It is important to ensure equitable outcomes and avoid perpetuating or exacerbating existing biases in society.

53. What are some challenges associated with deploying machine learning models in real-world applications?

Challenges associated with deploying machine learning models in real-world applications include data quality and bias, interpretability and transparency, scalability and efficiency, robustness to adversarial attacks, ethical considerations, and regulatory compliance.

54. What is the role of interpretability in machine learning models?

Interpretability in machine learning models refers to the ability to understand and explain the underlying mechanisms and reasoning behind the model's predictions or decisions. It is essential for building trust, debugging models, and ensuring accountability in AI systems.

55. How can interpretability be achieved in machine learning models?

Interpretability in machine learning models can be achieved through various techniques such as using simpler and more transparent models, feature importance analysis, model-agnostic interpretability methods (e.g., SHAP, LIME), and providing explanations in natural language or visual form.

56. What are some ethical considerations in AI and machine learning?

Ethical considerations in AI and machine learning include issues related to fairness and bias, privacy and data protection, accountability and transparency, safety and security, societal impact and inequality, and human control and autonomy.

57. How can bias in AI and machine learning models be addressed?

Bias in AI and machine learning models can be addressed through various approaches such as diverse and representative data collection, bias detection and mitigation techniques, fairness-aware algorithm design, and interdisciplinary collaboration involving ethicists, domain experts, and affected communities.

58. What is data augmentation, and why is it used in machine learning?

Data augmentation is a technique used to artificially increase the size of a training dataset by applying transformations such as rotation, scaling, cropping, or adding noise to the existing data samples. It helps improve model generalization and robustness, especially when training data is limited.

59. What is the importance of data preprocessing in machine learning?

Data preprocessing is crucial in machine learning as it involves cleaning, transforming, and organizing raw data into a suitable format for training machine learning models. Proper data preprocessing can improve model performance, reduce overfitting, and ensure the quality and integrity of the data.

60. What are some common data preprocessing techniques?

Common data preprocessing techniques include handling missing values (imputation), feature scaling and normalization, encoding categorical variables, handling outliers, dimensionality reduction (e.g., PCA), and splitting the data into training, validation, and test sets.

61. What is imbalanced data, and how can it affect machine learning models?

Imbalanced data refers to datasets where the distribution of classes or labels is heavily skewed, with one or more classes being significantly more prevalent than others. Imbalanced data can lead to biased models, poor generalization, and low predictive performance, especially for minority classes.

62. What are some techniques for dealing with imbalanced data?

Techniques for dealing with imbalanced data include resampling methods such as oversampling (e.g., SMOTE) and undersampling, using different evaluation metrics (e.g., precision-recall curve, F1-score) instead of accuracy, and using ensemble methods that handle class imbalance effectively.

63. What is anomaly detection, and what are some common approaches to anomaly detection?

Anomaly detection is the task of identifying rare or unusual patterns or instances in data that deviate significantly from the norm. Common approaches include statistical methods, machine learning techniques (e.g., isolation forests, one-class SVM), and deep learning methods (e.g., autoencoders).

64. What is clustering, and what are some common clustering algorithms?

Clustering is the task of grouping similar objects or data points into clusters or segments based on their intrinsic properties or features.

Common clustering algorithms include k-means clustering, hierarchical clustering, DBSCAN, and Gaussian mixture models.

65. What is the difference between supervised and unsupervised learning?

Supervised learning involves training a model on labeled data, where the input features are associated with corresponding target labels or outcomes. Unsupervised learning involves training a model on unlabeled data, where the model learns to find patterns or structures in the data without explicit guidance.

66. What is the curse of dimensionality, and how does it affect machine learning models?

The curse of dimensionality refers to the phenomena where the performance of machine learning models deteriorates as the number of features or dimensions in the data increases. It leads to increased computational complexity, sparsity of data, and decreased generalization ability of models.

67. What is dimensionality reduction, and why is it used in machine learning?

Dimensionality reduction is the process of reducing the number of input features or dimensions in a dataset while preserving the most important information. It is used in machine learning to address the curse of dimensionality, improve model performance, and visualization of high-dimensional data.

68. What are some common dimensionality reduction techniques?

Common dimensionality reduction techniques include principal component analysis (PCA), t-distributed stochastic neighbor embedding (t-SNE), linear discriminant analysis (LDA), and autoencoders.

69. What is principal component analysis (PCA), and how does it work?

Principal component analysis (PCA) is a popular dimensionality reduction technique that transforms high-dimensional data into a lower-dimensional representation by projecting it onto a new coordinate system defined by the principal components. These components are orthogonal vectors that capture the maximum variance in the data.

70. What is natural language processing (NLP)?

Natural language processing (NLP) is a field of artificial intelligence that focuses on enabling computers to understand, interpret, and generate human language. It involves tasks such as text classification, sentiment analysis, named entity recognition, machine translation, and text generation.

71. What are some common applications of natural language processing (NLP)?

Common applications of natural language processing (NLP) include sentiment analysis, chatbots and virtual assistants, machine translation, text summarization, named entity recognition, question answering systems, and information retrieval.

72. What is tokenization in natural language processing (NLP)?

Tokenization is the process of breaking a text into smaller units called tokens, which can be words, phrases, or symbols. It is a fundamental preprocessing step in NLP tasks such as text classification, named entity recognition, and machine translation.

73. What is word embedding, and how is it used in natural language processing (NLP)?

Word embedding is a technique used to represent words as dense, low-dimensional vectors in a continuous vector space, where semantically similar words are mapped to nearby points. It is used in NLP to capture semantic relationships between words, improve model performance, and enable word-level analysis.

74. What are some common word embedding techniques?

Common word embedding techniques include Word2Vec, GloVe (Global Vectors for Word Representation), fastText, and BERT (Bidirectional Encoder Representations from Transformers).

75. What is Word2Vec, and how does it work?

Word2Vec is a popular word embedding technique that learns distributed representations of words in a continuous vector space from large corpora of text data. It uses shallow neural networks to predict the context words given a target word (skip-gram model) or predict the target word given a context (continuous bag-of-words model).

76. What is sentiment analysis, and how is it performed?

Sentiment analysis is the task of determining the sentiment or emotional tone expressed in a piece of text, such as positive, negative, or neutral. It is performed using machine learning techniques, where the text is represented as features, and a model is trained to classify the sentiment of the text.

77. What are some challenges in sentiment analysis?

Challenges in sentiment analysis include handling sarcasm, irony, and ambiguity in language, dealing with context-dependent sentiments, addressing domain-specific language and slang, and ensuring the accuracy and reliability of sentiment labels in training data.

78. What is named entity recognition (NER), and why is it important in natural language processing (NLP)?

Named entity recognition (NER) is the task of identifying and classifying named entities (such as names of persons, organizations, locations, dates, etc.) mentioned in text data. It is important in NLP for information extraction, document summarization, and question answering.

79. What is the difference between a generative model and a discriminative model?

A generative model learns the joint probability distribution of the input features and the target labels, enabling it to generate new samples similar to the training data. A discriminative model learns the conditional probability distribution of the target labels given the input features and is primarily used for classification tasks.

80. What is a Markov chain, and how is it used in natural language processing (NLP)?

    A Markov chain is a stochastic model that describes a sequence of events where the probability of each event depends only on the state of the preceding event. In NLP, Markov chains are used for tasks such as text generation, speech recognition, and part-of-speech tagging.

81. What is topic modeling, and what are some common techniques for topic modeling?

    Topic modeling is a technique used to discover latent topics or themes present in a collection of documents. Common techniques for topic modeling include Latent Dirichlet Allocation (LDA), Latent Semantic Analysis (LSA), and Non-negative Matrix Factorization (NMF).

82. What is Latent Dirichlet Allocation (LDA), and how does it work?

    Latent Dirichlet Allocation (LDA) is a probabilistic generative model used for topic modeling. It represents each document as a mixture of latent topics, where each topic is characterized by a distribution over words. LDA infers the latent topics from the observed words in the documents.

83. What is Latent Semantic Analysis (LSA), and how does it work?

    Latent Semantic Analysis (LSA) is a technique used for dimensionality reduction and semantic analysis of text data. It represents documents and terms as vectors in a high-dimensional space and applies singular value decomposition (SVD) to identify latent semantic relationships between words and documents.

84. What is Non-negative Matrix Factorization (NMF), and how does it work?

    Non-negative Matrix Factorization (NMF) is a dimensionality reduction technique used for clustering and topic modeling of non-negative data, such as text documents or images. It factorizes the input matrix into two lower-dimensional matrices, ensuring that all elements are non-negative.

85. What is machine translation, and how does it work?

Machine translation is the task of automatically translating text from one language to another using computer algorithms. It works by training machine learning models (e.g., neural machine translation models) on parallel corpora of translated sentences, learning to generate translations from input text.

86. What are some challenges in machine translation?

Challenges in machine translation include handling linguistic variations and nuances across languages, preserving meaning and context during translation, dealing with idiomatic expressions and cultural differences, and achieving high translation accuracy and fluency.

87. What is sequence-to-sequence learning, and how is it used in machine translation?

Sequence-to-sequence learning is a neural network architecture used for tasks where the input and output are sequences of arbitrary lengths, such as machine translation. It consists of an encoder-decoder framework, where the encoder processes the input sequence into a fixed-size representation, and the decoder generates the output sequence based on this representation.

88. What is attention mechanism, and how is it used in sequence-to-sequence models?

Attention mechanism is a technique used in sequence-to-sequence models to improve the quality of generated sequences by focusing on relevant parts of the input sequence during decoding. It allows the model to dynamically weigh the importance of different input tokens based on their relevance to the current decoding step.

89. What is speech recognition, and how does it work?

Speech recognition is the task of automatically transcribing spoken language into text. It works by processing audio signals using machine learning algorithms (e.g., deep neural networks) to extract features, such as spectrograms or MFCCs, and mapping them to text sequences.

90. What are some challenges in speech recognition?

Challenges in speech recognition include dealing with variations in speech patterns and accents, background noise and environmental conditions, speaker diarization and speaker adaptation, handling out-of-vocabulary words, and achieving real-time performance and low latency.

91. What is reinforcement learning, and how does it differ from supervised learning and unsupervised learning?

    Reinforcement learning is a type of machine learning where an agent learns to make decisions by interacting with an environment to maximize cumulative rewards. Unlike supervised learning, reinforcement learning does not require labeled data, and unlike unsupervised learning, it involves learning from feedback (rewards) received from the environment.

92. What are some applications of reinforcement learning?

    Applications of reinforcement learning include game playing (e.g., AlphaGo), robotic control and automation, recommendation systems, financial trading, autonomous vehicles, and healthcare optimization.

93. What is the difference between model-based and model-free reinforcement learning?

    In model-based reinforcement learning, the agent learns a model of the environment's dynamics (transition function and reward function) and uses this model to plan its actions. In model-free reinforcement learning, the agent directly learns a policy or value function from interactions with the environment without explicitly modeling its dynamics.

94. What is the exploration-exploitation tradeoff in reinforcement learning?

    The exploration-exploitation tradeoff in reinforcement learning refers to the dilemma faced by an agent when deciding whether to explore new actions or exploit the current knowledge to maximize rewards. Balancing exploration and exploitation is essential for discovering optimal policies in reinforcement learning tasks.

95. What are some exploration strategies in reinforcement learning?

Exploration strategies in reinforcement learning include ε-greedy exploration, softmax exploration, Upper Confidence Bound (UCB) exploration, Thompson sampling, and exploration based on intrinsic motivation or curiosity.

96. What is deep reinforcement learning, and how does it differ from traditional reinforcement learning?

Deep reinforcement learning is a subfield of reinforcement learning that combines deep learning techniques with reinforcement learning algorithms to handle high-dimensional input spaces and complex decision-making tasks. It differs from traditional reinforcement learning by using deep neural networks to approximate value functions or policies.

97. What are some challenges in deep reinforcement learning?

Challenges in deep reinforcement learning include sample efficiency (i.e., learning from limited interactions with the environment), instability and convergence issues in training deep neural networks, exploration in high-dimensional action spaces, and generalization to unseen environments.

98. What is the role of reward shaping in reinforcement learning?

Reward shaping is a technique used in reinforcement learning to design additional reward functions that provide informative feedback to the agent, guiding it towards desired behaviors and accelerating learning. Reward shaping can improve sample efficiency and convergence speed in reinforcement learning tasks.

99. What are some techniques for addressing the exploration-exploitation tradeoff in reinforcement learning?

Techniques for addressing the exploration-exploitation tradeoff in reinforcement learning include ε-greedy exploration, softmax exploration, Upper Confidence Bound (UCB) exploration, Thompson sampling, and intrinsic motivation or curiosity-driven exploration.

100. What are some ethical considerations in the deployment of reinforcement learning systems?

Ethical considerations in the deployment of reinforcement learning systems include concerns related to safety and risk management, fairness and bias in decision-making, accountability and transparency of algorithms, privacy and data protection, and societal impact and human welfare. It is essential to consider these ethical implications and ensure responsible and ethical use of reinforcement learning technology.

Data Science - https://youtube.com/live/C0LVx2QQYbs?feature=share

IoT - https://youtube.com/live/mNmg0AHBVM8?feature=share

Full stack Development - https://youtube.com/live/periQ_9kZWA?feature=share

---

Learn for Free

1. HTML	html.com
2. CSS	web.dev/learn/css
3. JavaScript	javascript.info
4. React	reactplay.io
5. Vue	learnvue.co
6. Git	git-scm.com/book
7. Web3	learnweb3.io
8. Python	learnpython.org
9. SQL	w3schools.com/sql
10. Blockchain	cryptozombies.io
11. Nextjs	nextjs.org/learn
12. AI	elementsofai.com
13. PHP	phptherightway.com
14. API	rapidapi.com/learn
15. GO	learn-golang.org